
Towards Perfect Code Generation: Iterative API-based Feedback Loops

Harman Preet Singh

852863316

Large language models are good at writing HTML that works, but not always HTML that is semantically correct. A common failure mode is ‘divitis’: wrapping everything in `<div>` tags instead of using semantic elements like `<nav>` or `<article>`, which makes pages harder to maintain and harder for screen readers to interpret. This report tests whether an iterative feedback loop, where an external validator flags structural errors and the model is reprompted with that feedback, can reduce structural errors in HTML output. We find that the feedback loop significantly reduces validation errors, with a 76.3% mean error reduction in the feedback condition. Additionally, a blind ablation check confirms that this improvement is attributable to the validator’s specific error feedback, not to reprompting alone. We also find that smaller models with native reasoning capability can match or exceed the output quality of models with far more parameters, and that the guardrail enables this parity without increasing token cost. Finally, we determine that a maximum of 2 reprompt iterations captures the majority of resolvable errors (99% of converging runs do so by iteration 2), beyond which diminishing returns make further iterations wasteful.

Keywords: semantic HTML · LLM code generation · iterative feedback loops

1 Introduction

Large Language Models (LLMs) are increasingly used in the field of software development for the generation of code, including HTML. These models are very capable of producing syntactically valid HTML that renders correctly in web browsers. However, there is a significant blind spot: while the generated code may be syntactically correct, it often lacks semantic correctness. In other words, the code functions as intended, but it may not adhere to best practices.

This potential non-adherence to semantic best practices is often seen in the overuse of the `<div>` tag, a phenomenon colloquially referred to as "divitis". Instead of using semantically meaningful tags in layout and content structure, such as `<section>` to denote sections or `<article>` for articles, developers may default to using `<div>` tags. For example, a developer might wrap an entire article in a `<div>` instead of using the more appropriate `<article>` tag. LLMs are trained on real code snippets, which often contain this pattern, making them likely to reproduce it in their generated output. Additionally, if a codebase already contains a significant amount of "divitis" code-generation tools like Claude Code may produce new code that perpetuates this issue.

There are two real consequences. Code that is hard to maintain, as the semantic meaning of the content has become an afterthought. But more importantly, users with visual impairments who rely on screen readers are negatively affected because screen readers use the semantic structure of a webpage to navigate them. A lack of semantic structure can make it difficult for these users to understand the content and context of a page.

Instructing an LLM to use semantic tags through a system prompt or explicit instructions within the prompt itself does not guarantee compliance, particularly for smaller or local models. There is a need for a reliable automated mechanism that can catch and correct this class of semantic errors in LLM-generated HTML.

We propose a solution in the form of an iterative validation-and-regeneration feedback loop. In this approach, the LLM-generated HTML is first validated using a validator tool which can flag semantic issues. The LLM is then reprompted with the identified errors, allowing it to correct its output. This loop continues until the output is free of errors or a maximum number of iterations is reached.

In this report, we aim to find out (RQ1) to what extent an iterative feedback loop can improve the semantic correctness of LLM-generated HTML code, and (RQ2) whether this guardrail approach can enable smaller,

cheaper LLMs to achieve semantic correctness comparable to larger models. The remainder of this report is structured as follows: Section 2 reviews related work, Section 3 presents the research questions and hypotheses, Section 4 details the methodology, Section 5 describes the experimental setup, Section 6 presents the results, and Section 7 discusses the findings and implications. Finally, we conclude with a summary of our contributions and potential future work.

2 Related Work

Related work spans two interconnected areas: (1) LLM code generation quality, and (2) iterative self-correction and feedback loops. In the first part, we discuss what the literature says about LLMs generating code: their strengths, but also the documented tendency to reproduce patterns from training data including bad practices. In the second part, we cover approaches where LLMs are given feedback and asked to revise their output, particularly for Answer Set Programming (ASP) code. Finally, we review what is known about HTML semantic quality as a measurable property and the use of automated validation as a guardrail.

The literature on LLM code generation highlights both the strengths and limitations of these models. While LLMs can autonomously generate functioning code, they are also prone to reproducing patterns from their training data, including suboptimal or even harmful coding practices. For instance, Souly et al. [1] demonstrate that LLMs can be susceptible to poisoning attacks, where maliciously crafted inputs lead to the generation of flawed code. They show that even a small amount of poisoned data can significantly degrade the model’s performance.

Tambon et al. [2] examined samples of 333 bugs collected from code generated using three leading LLMs and identified 10 common bug patterns, not including our focus on semantic quality. Dou et al. [3] evaluated several models on the length and complexity of generated code, finding that LLMs face challenges in generating successful code for more complex problems. They propose a novel training-free iterative method where the model gains the ability to self-critique and correct its own code, resulting in a repair success rate of 29.2% after two iterations. This study particularly is very relevant to our work, as it demonstrates the potential of iterative feedback loops in improving code quality. However, it primarily focuses on functional correctness rather than semantic or stylistic correctness, which is a gap that our research aims to address.

Addressing the second area of related work, several studies have explored iterative feedback and self-correction loops for LLMs. Ravi et al. [4] present a method that employs multiple feedback models to iteratively refine code generated by LLMs. Their approach demonstrates that iterative feedback can significantly enhance code quality, but it also introduces potential biases from the feedback models themselves. In contrast, our approach does not use other LLMs as feedback sources; instead, we employ an external, deterministic validator to provide feedback. This choice mitigates the risk of introducing additional biases and allows for a more objective assessment of code quality.

Madaan et al. [5] introduce ‘Self-Refine’, a self-refinement framework where the model critiques its own output and iteratively improves it. They tested the framework out on 7 diverse tasks, ranging from dialog response generation to mathematical reasoning, and showed that the generated code improves by 20% on average in task performance. The proposed method uses a single LLM as the generator, refiner and the feedback provider, whereas our approach separates the feedback mechanism from the generation process.

Much research has been conducted comparing the performance of smaller versus larger LLMs on code generation tasks. Coignon et al. [6] show that larger models generally outperform smaller ones in terms of code quality and functional correctness. However, they note that actual runtime execution efficiency of the generated code does not vastly differ, that is, smaller models can still produce code that runs efficiently.

Historically, treating semantic correctness as a measurable, objective property has been difficult because ‘meaning’ is highly contextual. However, some prior work has attempted to quantify semantic correctness in code. One example is the ‘Generic-to-Semantic Ratio’ (R_{gs}), which measures the count of non-semantic structural nodes (`<div>`, `<span>`) against native layout elements (`<header>`, `<nav>`, `<main>`, `<section>`, `<footer>`). A lower ratio indicates higher structural semantic quality. The ‘Heading Hierarchy Traversal’ (H_{score}) method measures structural correctness by scanning a webpage for skipped heading levels (e.g., a jump from `<h1>` directly to `<h3>`, skipping `<h2>`). These approaches provide a way to quantify semantic quality, but they are limited in scope and do not capture the full semantic intent of the code. In comparison, the W3C Markup Validation Service is a widely used tool that checks for syntactic and structural conformance to HTML standards. It is capable of identifying errors such as unclosed tags, incorrect nesting, and invalid attributes. However, it does not assess the semantic intent of the code, which is a critical aspect of web accessibility and user experience. This gap highlights the need for more comprehensive evaluation methods that can capture both syntactic correctness and semantic quality in HTML code. Our work does not address this gap, instead we acknowledge the need for further research in this area.

No prior work combines an external deterministic validator, an iterative reprompting loop, and a cost-efficiency comparison across model sizes in the context of HTML semantic quality. Our research aims to fill this gap by

providing a comprehensive evaluation of how iterative feedback loops can improve the semantic quality of code generated by LLMs. By using an external validator, we can objectively assess the quality of the generated code without introducing additional biases. Additionally, our study will explore the cost-efficiency of using smaller models in comparison to larger ones.

3 Research Questions & Hypotheses

We aim to answer two research questions:

1. **RQ1:** Does the iterative API-based feedback loop improve the quality of code generation? We operationalise "quality" as the reduction in the number of errors in the generated code, measured before and after applying the feedback loop.
2. **RQ2:** Does this guardrail approach enable smaller, cheaper LLMs to achieve comparable quality to larger, more expensive models? We define "comparable quality" as equivalent error reduction rates, and "cost" is measured via mean generation time per prompt per model. Cloud-based models are explicitly out of scope for this study, as all models are run locally through Ollama, ensuring consistent evaluation conditions and avoiding external dependencies.

To evaluate our feedback loop, we also test the models under a blind condition, where the regeneration is performed without detailed feedback. This allows us to isolate the effect of the feedback loop.

We expect the feedback condition to outperform the blind condition across all models, as the iterative feedback loop provides targeted guidance for error correction. Furthermore, we hypothesise that reasoning-capable models (i.e., models that think and reason through the problem before generating code) will converge faster regardless of parameter count. Additionally, we expect that difficult prompts will require more iterations to achieve error-free code generation.

4 Methodology

4.1 System Architecture

The system consists of three key components. First, we have the language models themselves, which are served locally and ran via Ollama. The second and most important component, the validator, is the Nu Html Checker (vnu) service. This service is available via a public API and also, conveniently, as a Docker container ¹. Finally, the third component is the Python pipeline that handles the connection between the LLM and the validator.

Each of these components is its own distinct module. This modularity is a deliberate design choice, allowing for the LLM and validator to be swapped out without needing to rewrite the pipeline. We run the validator locally in Docker rather than using the public API for two reasons. First, the public API has rate limits that would make large-scale batch runs across 51 prompts, multiple models, and multiple conditions impractical. Second, running the validator locally allows us to pin the version and ensure consistency across all runs. This is important to ensure reproducibility of our results.

The pipeline, which is illustrated in Figure 1, operates in a loop. First, the prompt is sent to the LLM, which generates the initial HTML code. This code is validated by the validator. If any issues are found, the pipeline reprompts the LLM with the original code and the list of issues. This loop continues until either the code passes validation or the maximum number of iterations is reached.

4.2 Prompt Dataset

The prompt dataset was hand-crafted rather than sampled from an existing benchmark, because no such benchmark exists for HTML semantic quality specifically. The goal was to create tasks that vary in difficulty structurally. The full prompt dataset is included in the Appendix.

We organise the prompts into a three-tiered structure: simple, medium, and difficult. The difficulty of a prompt is based on the structural complexity of the desired output. Simple prompts involve single components, such as a couple headers that contain some text. Medium prompts involve multi-section layouts, such as a navigation bar, a main content area, and a footer. Difficult prompts involve complex structures, such as the periodic table of elements, which requires a complex table layout with many containers to fit the name of the chemical element, its symbol, its atomic number, and other information. The more complex the structure, the more opportunities there are for the model to make a semantic error.

¹ ghcr.io/validator/validator:latest




Figure 1: Overview of the iterative validation and regeneration loop.

All models are evaluated on the exact same 51 prompts in the same order. We opted for this fixed design to ensure that differences in results are attributable to the model or condition, rather than to prompt sampling variance.

4.3 Models & Configurations

We selected a range of models that vary in both parameter count and reasoning capability. Additionally, all models are available through Ollama and run locally at full precision (no quantisation was used), which keeps the hardware environment constant. An interesting finding from our initial testing was that reasoning capability matters more than parameter count. Smaller models with reasoning capability outperformed larger models without it. This finding shaped our final model selection, as we wanted to ensure that both axes (parameter count and reasoning capability) were represented in our experiments. An overview of all evaluated models is provided in Table 1. The table includes the model name, parameter count, reasoning capability (yes/no)

Model	Parameters (B)	Reasoning
gemma3:1b	1	No
qwen3:4b-thinking	4	Yes
deepseek-coder:6.7b	6.7	No
gemma4:e4b	8	Yes
qwen2.5-coder:14b	14.8	No
codestral:22b	22	No

Table 1: Overview of all evaluated models

We set the maximum iteration count as a parameter of interest. For each (model, prompt, condition) combination, we run the pipeline up to the maximum number of iterations, recording when convergence occurs. This allows us to analyse diminishing returns and determine the practical cut-off beyond which additional iterations yield negligible improvements. This is important for understanding the cost-quality tradeoff in RQ2, as it tells us how many iterations are really necessary.

4.4 Validation Strategy

The validator checks structural and syntactic conformance to the HTML specification. A limitation of the validator is that it cannot evaluate semantic intent. For example, it cannot tell you that `<div class="nav">` is worse than `<nav>` because both are syntactically valid.

Despite this limitation, the validator is still a meaningful signal. It can still detect many structural issues such as improper nesting, missing landmark elements, and attribute misuse. While error/warning count reduction is a necessary but not sufficient proxy for semantic improvement, it is still useful and can be triangulated with other measures.

In addition to running the pipeline in the feedback condition, we also run it in a blind condition. In the blind condition, the model is simply told to review and improve its output, with no error details provided. This allows us to isolate the contribution of the validator feedback itself, distinguishing it from the general benefit of giving a model a second attempt.

In related work where code quality was evaluated, the validator used was other LLMs, which struggle to provide deterministic outputs. To ensure that the validator is deterministic, we run a check on the same HTML input multiple times to verify that the validator produces identical results. This is non-trivial. If the validator were stochastic, our improvement metrics would be meaningless. We found that the validator is indeed deterministic, which gives us confidence in our before/after deltas.

4.5 Cost Metric

Since all models run locally via Ollama, there is no API pricing to measure. Instead, we proxy cost through three values captured per iteration: `gen_time_s` (wall-clock generation time), `prompt_eval_count` (tokens in the prompt), and `eval_count` (tokens generated). These together give a picture of both time cost and computational load per run.

We acknowledge that no cloud-hosted large model (e.g., GPT-4o, Claude) is included in this study. The comparison is therefore entirely within the local model ecosystem. This is a deliberate scope boundary, as cloud API costs introduce pricing variables outside our control. We note this as a direction for future work.

5 Experimental Setup

As mentioned in subsection 4.3, prior to the full experimental run, a small-scale pilot was conducted on a subset of prompts using a small variety of models. The most striking observation from this pilot was that models with native reasoning capabilities consistently resolved all validator errors within a single iteration, whereas non-reasoning models did not. Such models often were not able to remove all errors within the pre-defined iteration limit. This finding directly informed the decision to treat reasoning capability as an explicit variable in the final design, rather than focusing solely on parameter count. Furthermore, the pilot revealed a practical issue with local model outputs: models wrapped their HTML outputs in markdown code fences due to how these models are trained to generate outputs. Occasionally, the models would prefix responses with explanatory prose, which the validator would reject as incorrect HTML syntax. For this reason, for the full experimental run we introduce a module which strips such formatting artefacts before validation so that presentation-level slip-ups are not conflated with HTML errors. To conclude, the pilot served a dual purpose: it validated the core loop design and revealed the preprocessing steps necessary for a fair before/after comparison

6 Results

6.1 Guardrail Effectiveness (RQ1)

The error counts reported here are derived from the validator’s output, which is a proxy for semantic correctness rather than a direct measure of it. As previously mentioned, while the validator provides a useful automated assessment of HTML validity, it may not capture all aspects of semantic correctness. We report error and warning counts as our primary metrics, but the results should be interpreted with this caveat in mind.

6.1.1 Overall before/after results (feedback condition)

Overall, across all models and prompts in the feedback condition ($n = 553$ runs), we observed a significant reduction in both errors and warnings after applying the guardrail. The aggregate results are summarized in Table 2. The mean error count decreased from 17.16 to 4.07, representing a 76.3% reduction, while the mean warning count decreased from 6.00 to 1.49, representing a 75.2% reduction. Paired Wilcoxon signed-rank tests confirmed that these reductions were statistically significant ($p < 0.001$), with 95% confidence intervals indicating meaningful decreases in both metrics.

Metric	Mean before	Mean after	Δ	% reduction	p	95% CI (Δ)
Errors	17.16	4.07	13.10	76.3%	< 0.001	[4.16, 22.03]
Warnings	6.00	1.49	4.51	75.2%	< 0.001	[0.01, 9.02]

Table 2: Aggregate before/after comparison

6.1.2 Feedback vs. Blind Ablation

In Table 3, we present the mean error reduction (before – after) by prompt difficulty tier and condition, split by whether the model has a native reasoning/thinking mode. Higher values indicate more errors were resolved. Notably, in each difficulty tier, the feedback condition consistently outperforms the blind condition for both reasoning and non-reasoning models. This suggests that the validator’s specific error list is indeed contributing to meaningful improvements in code generation quality, beyond just reprompting.

Difficulty	Condition	Reasoning models	Non-reasoning models
Simple	Feedback	0.33 (n=30)	0.18 (n=67)
Simple	Blind	0.13 (n=30)	-0.21 (n=67)
Medium	Feedback	0.55 (n=73)	0.38 (n=141)
Medium	Blind	0.14 (n=73)	-0.27 (n=141)
Difficult	Feedback	1.86 (n=28)	-0.10 (n=61)
Difficult	Blind	-1.25 (n=28)	-0.46 (n=61)

Table 3: Mean error reduction (before – after)

Across all difficulty tiers, the feedback condition shows a clear advantage over the blind condition, particularly for reasoning models. For example, in the difficult tier, reasoning models in the feedback condition resolved an average of 1.86 errors, while those in the blind condition actually saw an increase in errors (-1.25). Across simple and medium tiers, the blind condition does not result in an increase in errors, but it still underperforms the feedback condition. This suggests that the specific feedback provided by the validator is useful for guiding the model towards more correct code generation, especially for models with reasoning capabilities.

In Figure 2, we visualise the mean error count per reprompt iteration for each model, comparing the feedback and blind conditions. The shaded bands represent ± 1 standard error of the mean (SEM). The symlog y -axis is used to accommodate the wide range of error counts across models, which span roughly 0–60 errors. We observe that the feedback condition consistently decreases the mean error count across iterations, as does the blind condition. However, the feedback condition shows a steeper decline in errors, ending 67% lower on average. Multiple models in the feedback condition reach near-zero error counts by the second or third iteration, while the blind condition never does. The two best performing models for both conditions are the reasoning model `qwen3:4b-thinking` and the non-reasoning model `codestral:22b`, the latter of which has the largest parameter count, but is outperformed by the smaller reasoning model.




Figure 2: Mean error count per reprompt iteration

6.1.3 Error Reduction by Prompt Difficulty Tier

In Figure 3, we see that difficult prompts have higher initial error counts than medium or simple prompts, with a mean of 14.01 errors for the initial condition across all models, compared to 0.48 for simple prompts and 10.80 for medium-difficulty prompts. The feedback condition lowers this mean error count of 14.01 to a mean of 4.87, while the blind condition only reduces it to 6.57. These are reductions of ~65% and ~53%, respectively, which show that the feedback condition outperforms the blind condition. Surprisingly, simple prompts ended up with a higher mean final error count for both the feedback and blind conditions than the initial error count. This is likely due to the fact that simple prompts have fewer lines of code, so any single error has a larger impact on the overall error count. So if a model introduces a new error during reprompting, it can outweigh the resolution of existing errors, leading to an overall increase in the mean error count.




Figure 3: Mean final error count (95% CI) by prompt difficulty tier and condition

6.1.4 Error Category Breakdown

The validator is capable of detecting a wide range of errors, which allows us to analyse which errors are most effectively resolved by the guardrail. Table 4 shows the frequency of each error category before and after applying the guardrail, across all models and runs. We see that the most common error category is "other," which is a catch-all for uncategorised errors. Looking at the errors that are categorised, we see that "disallowed-attribute" errors are the most effectively resolved, with a 96% reduction. This is likely due to the fact that disallowed attributes are often easy to identify and remove, and the guardrail can provide specific guidance on which attributes are not allowed. On the other hand, "disallowed-child" errors are less effectively resolved, with only a 28% reduction.

There are three categories where the number of errors actually increased after applying the guardrail: 'missing-doctype', 'cannot-recover' and 'missing-lang'. Compared to the other categories, these errors have relatively low counts, so the increase in their frequency is not very statistically significant.

Category	Before	After	Δ	% resolved
other	11369	2619	8750	77.0%
disallowed-attribute	697	28	669	96.0%
stray-tag	324	100	224	69.1%
disallowed-child	314	225	89	28.3%
missing-title	61	32	29	47.5%
missing-doctype	18	29	-11	-61.1%
duplicate-tag	14	13	1	7.1%
cannot-recover	10	11	-1	-10.0%
missing-attribute	5	3	2	40.0%
missing-lang	0	1	-1	—

Table 4: Frequency of each validator message category

6.2 Cost-Efficiency (RQ2)

Given that these experiments are conducted on local models, we define cost-efficiency in terms of operational metrics rather than monetary cost. Specifically, we measure cost as the total generation latency (in seconds) and the total token consumption (the sum of prompt evaluation count and evaluation count) per run. Quality is assessed based on the error reduction achieved by each model.

6.2.1 Quality vs. Cost per Model

Figure 4 shows a comparison between the total time a model took per run, and the amount of errors left after all iterations. Looking first at the models without a native thinking mode, we can see that there is a general trend between the amount of parameters and the mean total generation time per run. However, this increase in generation time does not translate to a better final error count (Figure 5). For example, `deepseek-coder: 6.7b` (6.7B parameters) achieves a lower mean final error count than `qwen2.5-coder: 14b` (14B parameters). This suggests that the parameter count alone is not a sufficient predictor of model performance in terms of error reduction.




Figure 4: Parameter count vs mean total generation time per run, coloured by native thinking mode.




Figure 5: Mean total generation time per run vs mean final error count (feedback condition)

6.2.2 Reasoning vs. Non-Reasoning as the Decisive Variable

Instead, the presence of a native reasoning mode appears to be a more significant factor in achieving better quality outputs. We can see this in the reasoning model `qwen3:4b-thinking` (4B parameters) which achieves a lower mean final error count than the non-reasoning model `codestral:22b` (22B parameters), despite having roughly a fifth of the parameters. The former’s reasoning capabilities allow it to outperform the latter in terms of error reduction, while also being more cost-efficient in terms of total generation time per run.




Figure 6: Parameter count vs mean final error count (feedback condition)

For Figure 6 in particular, it is important to note that the y-axis only spans from a mean final error count of 0 to ~0.7 (less than 1). The absolute difference between models is relatively small, which is why the y-axis is zoomed in to better illustrate the differences between models. However, even small differences in mean final error count can be significant in practice, especially when considering the cumulative effect across multiple runs or iterations.

6.2.3 Cost of the Feedback Loop vs. the Blind Condition

The feedback loop is expected to incur a higher token cost than the blind condition, because it includes the original code plus a specific error list in each reprompt iteration. Compared to the blind condition, which only includes the original code, the feedback condition should require more tokens per iteration. However, the actual token cost per run may vary depending on how quickly each condition converges to a solution. If the feedback condition resolves errors more efficiently, it may require fewer iterations overall, potentially offsetting the higher per-iteration token cost.

Model	Mean feedback tokens	Mean blind tokens	Extra tokens (<i>feedback-blind</i>)	Extra errors resolved	Tokens / extra resolved
All models (pooled)	9710	11293	-1584 (-14.0%)	-0.049	—
codestral:22b	1356	1998	-642 (-32.1%)	+0.333	-1925.4
deepseek-coder:6.7b	2703	3083	-380 (-12.3%)	+0.925	-410.6
gemma3:1b	2999	2528	+470 (+18.6%)	+0.132	3573.8
gemma4:e4b	8625	9172	-547 (-6.0%)	+3.100	-176.5
qwen2.5-coder:14b	3166	3855	-689 (-17.9%)	+0.545	-1262.8
qwen3:4b-thinking	6538	15584	-9046 (-58.0%)	+0.297	-30455.0
smollm2:135m	24292	23839	+453 (+1.9%)	-1.895	—

Table 5: Mean total tokens per run (initial generation plus all reprompt iterations) under feedback vs blind

In Table 5, we can see that the feedback condition is actually cheaper than the blind condition in total tokens per run for 5 of 7 models (and pooled across all models). This is because the blind condition converges more slowly (more reprompt iterations needed on average, each carrying its own token cost), so even though each individual feedback iteration’s prompt is longer (original code + specific error list), the feedback condition needs fewer of them.

The last column shows the number of tokens spent per additional error resolved, which is only defined where feedback resolves strictly more errors than blind. For example, `deepseek-coder:6.7b` spends an average of ~400 tokens per additional error resolved compared to blind, while `qwen3:4b-thinking` spends an average of ~30,000 tokens per additional error resolved. This suggests that while the feedback condition may be more efficient in terms of total token cost, the quality gain may not always justify the extra cost, especially for models that resolve fewer additional errors.

6.3 Iteration Count Analysis

75% of runs converge to zero errors within 3 iterations, and 99% of those runs do so by iteration 2, as shown in Figure 7. This suggests that the majority of resolvable errors are fixed within the first few iterations, and that additional iterations beyond this point yield diminishing returns. The distribution of convergence iterations is also influenced by prompt difficulty tier, with simple prompts converging more quickly than medium or difficult prompts. Interestingly, difficult prompts seem to converge more quickly than medium prompts, which also start with a higher mean error count.




Figure 7: Distribution of the reprompt iteration at which each (model, prompt, trial) run first reaches zero errors, feedback condition, split by prompt difficulty tier. “Not resolved” = still nonzero at the last observed iteration.

6.3.1 Diminishing Returns and Optimal Cut-off

The amount of issues resolved does not scale linearly with the number of iterations, as shown in Table 6. There is a clear diminishing return on the number of iterations, with the majority of resolvable errors being fixed within the first two iterations. For example, in the simple prompt tier, 70.6% of runs have resolved all errors by iteration 2, and this does not increase at all by iteration 3. In the medium prompt tier, 70.5% of runs have resolved all errors by iteration 2, and this only increases slightly to 70.8% by iteration 3. In the difficult prompt tier, 62.3% of runs have resolved all errors by the first iteration, and this only increases to 66.4% with two additional iterations. This suggests that a practical recommendation for the maximum iteration count is 2, as additional iterations beyond this point yield minimal improvements in error resolution.

Difficulty	Iteration cutoff	% resolved	Mean remaining	n runs
Simple	0	57.4%	0.97	136
Simple	1	67.6%	11.44	136
Simple	2	70.6%	0.80	136
Simple	3	70.6%	0.80	136
Medium	0	50.8%	20.76	295
Medium	1	61.7%	14.88	295
Medium	2	70.5%	9.85	295
Medium	3	70.8%	5.24	295
Difficult	0	49.2%	26.52	122
Difficult	1	62.3%	13.47	122
Difficult	2	66.4%	2.43	122
Difficult	3	66.4%	4.87	122

Table 6: Diminishing returns by prompt difficulty tier, showing the percentage of runs that have resolved all errors by each iteration cutoff, and the mean number of remaining errors for those that have not yet resolved.

7 Discussion

7.1 Interpretation of Results

Through our experiment we can confidently state that the iterative feedback loop significantly improves HTML validity, with a 76.3% reduction in mean errors and 75.2% in warnings, which are both statistically significant ($p < 0.001$). The feedback condition consistently and meaningfully outperforms the blind condition across all difficulty tiers, confirming that the validator’s specific error list is doing real work: it is not merely the act of reprompting, giving the model a second chance, that drives improvement.

The blind condition’s behaviour noteworthy. For reasoning models on difficult prompts, the blind condition actively worsens output (mean error reduction of -1.25). Without the targeted feedback the validator provides, the reasoning model "overthinks" its revision and introduces new errors.

The decisive factor in closing the gap between smaller and larger models is not parameter count but reasoning capability. The `qwen3:4b-thinking` model (4B parameters) achieves a lower mean final error count than `codestral:22b` (22B parameters) and does so more cost-efficiently. This directly answers RQ2: smaller models can match or exceed larger ones, but only when they possess native reasoning capability, not simply because the guardrail gives them more attempts.

In Table 5, we, counterintuitively, saw that the feedback condition is actually cheaper than the blind condition for 5 of 7 models in total token consumption, because it converges faster. The guardrail provides both a quality and cost advantage. However, the `qwen3:4b-thinking` model stood out as an outlier. It spends ~30,000 tokens per additional error resolved over the blind condition. This leads us to question whether the quality gain justifies the cost for reasoning models specifically.

Surprisingly, we found that simple prompts end with higher mean final error counts than their starting error counts (visible in Figure 3). This is because simple prompts produce short HTML outputs, so a single newly introduced error during reprompting has an outsized effect on the mean. This is an artefact of the metric rather than a genuine quality failure.

7.2 Validity of the Validation Approach

The validator measures syntactic and structural conformance, not semantic intent. A model that replaces every `<div>` with another `<div class="nav">` would pass the validator but solve nothing. Error reduction is therefore a necessary but not sufficient proxy for semantic correctness.

The validator does catch a meaningful class of structural errors that correlate with poor semantic quality even if they do not perfectly capture it. The 96% resolution rate for disallowed-attribute errors and 69% for stray-tag errors suggest the guardrail is correcting real, detectable problems.

In the related work we introduced the Generic-to-Semantic Ratio (Rgs) and Heading Hierarchy Score (Hscore) as metrics for semantic correctness. These were not applied in this study’s evaluation. Incorporating these metrics in future work would allow a more direct test of whether validator-driven error reduction actually translates to improved semantic structure.

7.3 Threats to Validity & Risk Analysis

The prompt dataset was hand-crafted, which introduces the risk that prompts were inadvertently designed in ways that favour certain model architectures or output styles. Medium-difficulty prompts behaving more erratically than difficult ones (Figure 3’s wider confidence intervals) may partly reflect this. The fixed 51-prompt ordering ensures all models face identical conditions, which limits within-experiment bias even if between-study generalisability is uncertain.

Even with a fixed seed and temperature setting, local model outputs are not guaranteed to be perfectly reproducible. We conducted repeated trials per (model, prompt, condition) combination to mitigate this, and the results were consistent across runs.

We implemented a validator determinism check to ensure that the validator’s output is consistent for the same HTML input. This was a non-trivial check, a stochastic validator would have made before/after deltas meaningless.

The findings are specific to HTML generation and locally available Ollama models. The core principle, which is that iterate validator feedback improved code quality and that reasoning capability matters more than parameter count, may generalise to other code generation contexts (SQL, Python, etc.) but this was not tested in this study. Future work should explore whether the same patterns hold for other programming languages and tasks. Additionally, the exclusion of cloud-based models (e.g., OpenAI, Anthropic) limits the applicability of the cost-efficiency findings, as these models have different pricing structures and performance characteristics.

8 Conclusion

The iterative feedback loop significantly reduces HTML validation errors, with a 76.3% mean error reduction in the feedback condition. The blind ablation confirms that this improvement is attributable to the validator’s specific error feedback, not to reprompting alone.

Smaller models with native reasoning capability can match or exceed the output quality of models with far more parameters. The guardrail enables this parity, and does so without increasing token cost. For most models, the feedback condition is actually cheaper in total tokens than the blind condition due to faster convergence.

A maximum of 2 reprompt iterations captures the majority of resolvable errors (99% of converging runs do so by iteration 2), beyond which diminishing returns make further iterations wasteful. A small reasoning-capable model under a feedback guardrail is a more cost-effective choice than a large non-reasoning model without one.

8.1 Future work

We determine three directions for future work. First, incorporate semantic metrics (Rgs, Hscore) alongside validator errors to measure divitis reduction. Second, extend the approach to include cloud-based models to test cost-efficiency at scale. And last, generalise the feedback loop to other code quality domains such as CSS accessibility or SQL schema correctness.

References

- [1] Alexandra Souly, Javier Rando, Ed Chapman, Xander Davies, Burak Hasircioglu, Ezzeldin Shereen, Carlos Mougán, Vasilios Mavroudis, Erik Jones, Chris Hicks, Nicholas Carlini, Yarin Gal, and Robert Kirk. Poisoning attacks on llms require a near-constant number of poison samples, 2025. URL <https://arxiv.org/abs/2510.07192>.
- [2] Florian Tambon, Arghavan Moradi-Dakhel, Amin Nikanjam, Foutse Khomh, Michel Desmarais, and Giuliano Antoniol. Bugs in large language models generated code: an empirical study. *Empirical Software Engineering*, 30, 02 2025. doi: 10.1007/s10664-025-10614-4.
- [3] Shihan Dou, Haoxiang Jia, Shenxi Wu, Huiyuan Zheng, Muling Wu, Yunbo Tao, Ming Zhang, Mingxu Chai, Jessica Fan, Zhiheng Xi, Rui Zheng, Yueming Wu, Ming Wen, Tao Gui, Xipeng Qiu, and Xuanjing Huang. What is wrong with your code generated by large language models? an extensive study. *Science China Information Sciences*, 69, 01 2026. doi: 10.1007/s11432-025-4632-8.
- [4] Ravin Ravi, Dylan Bradshaw, Stefano Ruberto, Gunel Jahangirova, and Valerio Terragni. Llmloop: Improving llm-generated code and tests through automated iterative feedback loops. 07 2025. doi: 10.1109/ICSME64153.2025.00109.
- [5] Aman Madaan, Niket Tandon, Prakhar Gupta, Skyler Hallinan, Luyu Gao, Sarah Wiegrefe, Uri Alon, Nouha Dziri, Shrimai Prabhumoye, Yiming Yang, Sean Welleck, Bodhisattwa Majumder, Shashank Gupta, Amir Yazdanbakhsh, and Peter Clark. Self-refine: Iterative refinement with self-feedback, 03 2023.
- [6] Tristan Coignon, Clément Quinton, and Romain Rouvoy. A performance study of llm-generated code on leetcode, 07 2024.

Appendix

Full prompt dataset

ID	Difficulty	Prompt
p01	Simple	Create a simple HTML page with a heading and a paragraph.
p02	Simple	Build an HTML page with a navigation bar containing three links: Home, About, Contact.
p03	Simple	Generate an HTML document with a table displaying 5 rows and 3 columns of sample data.
p04	Simple	Create a basic HTML form with fields for name, email, and a submit button.
p05	Simple	Write an HTML page that includes an image and a caption below it.
p06	Simple	Create an HTML layout with a header, main content area, and footer.
p07	Simple	Build an HTML page with an ordered and unordered list.
p08	Medium	Generate an HTML page with embedded CSS to style a button with hover effects.
p09	Medium	Create a responsive HTML page using a simple flexbox layout.
p10	Simple	Write an HTML document that includes a video player with controls.
p11	Medium	Create an HTML page with a login form styled using inline CSS.
p12	Medium	Build an HTML page with a sidebar and main content using CSS grid.
p13	Medium	Generate an HTML page that includes a dropdown menu navigation.
p14	Simple	Create an HTML document with semantic elements like <code><article></code> , <code><section></code> , and <code><aside></code> .
p15	Medium	Write an HTML page that displays a product card with image, title, price, and button.
p16	Medium	Create an HTML page with a contact form including validation attributes.
p17	Medium	Build an HTML page that uses Google Fonts and custom typography.
p18	Medium	Generate an HTML page with a dark mode toggle using basic JavaScript.
p19	Medium	Create an HTML layout mimicking a blog post page with comments section.
p20	Medium	Write an HTML page that includes a responsive image gallery.
p21	Medium	Create an HTML page with a fixed header and scrollable content.
p22	Simple	Build an HTML page that includes a progress bar styled with CSS.
p23	Medium	Generate an HTML page with tabs that switch content using JavaScript.
p24	Simple	Create an HTML form with radio buttons, checkboxes, and a dropdown.
p25	Simple	Write an HTML page that embeds a map using an iframe.
p26	Medium	Create an HTML page with a modal popup triggered by a button.
p27	Medium	Build an HTML page with a pricing table featuring three plans.
p28	Medium	Generate an HTML page with a sticky sidebar navigation.
p29	Medium	Create an HTML page with animations using CSS keyframes.
p30	Medium	Write an HTML page that includes a countdown timer using JavaScript.
p31	Difficult	Create an HTML dashboard layout with cards and charts placeholders.
p32	Difficult	Build an HTML page that mimics a social media post feed.
p33	Difficult	Generate an HTML page with a multi-step form wizard.
p34	Medium	Create an HTML page that uses ARIA roles for accessibility.
p35	Medium	Write an HTML page with a collapsible accordion component.
p36	Simple	Create an HTML page that includes breadcrumb navigation.
p37	Medium	Build an HTML page with a responsive navbar that collapses on mobile.
p38	Difficult	Generate an HTML page that displays a calendar layout.
p39	Medium	Create an HTML page with a file upload form and preview section.
p40	Medium	Write an HTML page that includes syntax-highlighted code blocks.
p41	Difficult	Create an HTML page with drag-and-drop functionality using JavaScript.
p42	Difficult	Build an HTML page that includes a carousel/slider component.
p43	Difficult	Generate an HTML page that mimics an e-commerce checkout flow.
p44	Difficult	Create an HTML page with a search bar that filters a list dynamically.
p45	Medium	Write an HTML page that includes a notification/toast system.
p46	Medium	Create an HTML page with a timeline component showing events.
p47	Medium	Build an HTML page with a rating system using stars.
p48	Difficult	Generate an HTML page that includes a chatbot UI interface.
p49	Medium	Create an HTML page that visualises hierarchical data using nested lists.
p50	Difficult	Write an HTML page that simulates a Kanban board with draggable cards.
p51	Difficult	Generate an HTML page that has the entire periodic table on it, formatted as a table.